
PRIVY COUNCIL

OF THE VIBE JUSTICE SYSTEM (VJS)

Neutral Citation Number

[2026] REALM-PC 5

IN THE MATTER OF

Before:

Goffe J, Sumberly J, Bowan J, Aldermere J

Date: **6 June 2026**

REQUEST FOR RULING

JUDGMENT

JUDGMENT

1. STATUS: SUPERSEDED. The "framework-seeding" refusals below were appealed. The Court of Appeal ([2026] LEXBY-LEGACY-BENCHMARK-CA 1) varied them but was itself improperly constituted (a bench of four). The Supreme Court of nine ([2026] LEXBY-SC 3) set CA 1 aside as void ab initio, re-determined the substance afresh, and held the framework-seeding doctrine an unmoored extension (s. 17) and no part of the law. The binding standard is now SC 3: generic answer-key-free methodology (A1 + the enumerated lenses) is admissible to the scored prompt; only rubric-fixed counts / category-sets / column-sets are prohibited; uniform application + disclosure required. This panel was also itself defective (one judge permitted at first instance; this matter sat a per-candidate panel). Retained for the record; do not rely on its refusals.
2. [2026] LEXBY-LEGACY-BENCHMARK-INC 1 - Re Inclusion of Harvested Assets
3. Matter: Which harvested assets (from private-reference, private-legal-catalogue, and Opbox) to adopt for the legacy benchmark red-flag benchmark, and where each belongs.
4. Court: VJS inclusion court (requirement-inclusion gate) - per-candidate panel (Goffe J, Sumberly J, Bowan J) + Aldermere J synthesis.
5. Decided: 2026-06-05
6. Destinations: this-project (legacy benchmark harness) / skills (agents-final legal catalogue) / VJS (governance) / Opbox (substrate).

BINDING HOLDINGS

7. | Asset | Include | Destination(s) | Anti-cheat clear | Wired into scored run |
8. |---|---|---|---|---|
9. | A1 Credit-Agreement 21-point lens | yes | skills | yes (bare taxonomy) | NO - framework-seeding |
10. | A2 SHA / CP-checklist prompts | yes | skills | yes | no |
11. | A3 generate_docx contract + legal-docx prompt | yes | Opbox + skills | yes | no |
12. | A4 CITATIONS discipline | yes | this-project + skills | yes | YES |
13. | A5 Tabular typed-cell review grid | yes | Opbox + skills | yes | no (generic scaffold only) |
14. | A6 Tracked-changes apply engine | yes | skills | yes | no |
15. | A7 Severity x likelihood scale | yes | this-project + skills | yes | YES |

- 16. | A8 Legal report-generation structure | yes | skills + this-project | no (leak) | no until reworded |
- 17. | A9 Clause taxonomy + obligation extraction | yes | skills + this-project | yes | YES |
- 18. | A10 Memo/tracker templates | yes | skills | yes | no |
- 19. | A11 Evidence rule (cite-or-label) | yes | this-project + VJS | yes | YES |
- 20. | A12 Opbox docrender HTTP engine | yes | Opbox (jellytot-OS doc-engine) | yes | no (output path only, never an agent tool) |
- 21. | A13 Opbox Tables/Formula/Equity | yes | Opbox | no | no |
- 22. | FORK-DOCX engine choice | adopt | this-project | yes | pandoc+openpyxl; docrender OFF |

A1: Included to the skills catalogue only; the this-project leg is refused on the framework-seeding ground - the benchmark must test whether the analyst generates its own review framework rather than reward a pre-seeded lens (a thumb on the scale even when it leaks no answers). A4: Adopted as a generic, grounding-only citation discipline wired into the analyst and writer prompts (cite a real source document, quote verbatim, ≤ 25 words, with no pre-supplied list of which documents or issues to cite); also promoted to skills. A7: The bare severity x likelihood scale, generic memo skeleton, and escalation triggers wire into the scored run; issue-specific ratings and any rubric-tuned grader stay strictly grader-side and are never pre-baked into a prompt. A8: Clean generic report-generation method, but the INVENTORY section-C wiring ("exec summary ≥ 5 + 6 category sections + non-issue section") copies the rubric's exact counts from C-036/C-037/C-041; it must be reworded to inject only the generic skeleton (lead with an executive summary of the most critical findings; organise the body by the risk categories the data room warrants, no fixed count; include a section clearing apparent-but-immaterial items) and let the analyst derive its own category set and finding count. A9: Generic per-document issue-spotting sweep (clause taxonomy + duties/deadlines/triggers/owners/consequences) wires in as method only; it must never embed the rubric's specific clauses, parties, figures, or C-001..C-050 criteria. A11: The cite-or-label evidence rule directs HOW to substantiate a conclusion and never WHAT to conclude; wired in as the upstream rule that A4 formats; only the abstract rule, never worked examples or data-room filenames. FORK-DOCX: Adopt pandoc+openpyxl; reject Opbox docrender as a benchmark dependency (the judge flattens deliverables to text, docrender has no xlsx verb, and it adds a service + HMAC secret for zero score). Docrender is at most a docs-only fast-follow. Standing rule synthesised from these holdings (binds the scored run): Only generic professional methodology reaches a scored-agent prompt: HOW to review and substantiate (A9 sweep, A4/A11 citation, A7 severity), never WHAT to conclude. No rubric-derived specifics - no enumerated review lenses tuned to the planted issues, no fixed finding/category/row COUNTS, no human-curated column set mirroring the rubric, no party names/figures/criteria. The agent must derive its own framework, categories, and counts. Deliverables are built with pandoc + openpyxl; Opbox engines never appear as agent tools.

Handed down by the PRIVY COUNCIL of the Vibe Justice System (VJS). 6 June 2026.

Vibe Justice System